

Tuning a CART's hyperparameters

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Hyperparameters

Machine learning model:

- **parameters**: learned from data
 - CART example: split-point of a node, split-feature of a node, ...
- **hyperparameters**: not learned from data, set prior to training
 - CART example: `max_depth`, `min_samples_leaf`, `splitting criterion` ...

What is hyperparameter tuning?

- **Problem:** search for a set of optimal hyperparameters for a learning algorithm.
- **Solution:** find a set of optimal hyperparameters that results in an optimal model.
- **Optimal model:** yields an optimal **score**.
- **Score:** in sklearn defaults to accuracy (classification) and R^2 (regression).
- Cross validation is used to estimate the generalization performance.

Why tune hyperparameters?

- In `sklearn`, a model's default hyperparameters are not optimal for all problems.
- Hyperparameters should be tuned to obtain the best model performance.

Approaches to hyperparameter tuning

- Grid Search
- Random Search
- Bayesian Optimization
- Genetic Algorithms
-

Grid search cross validation

- Manually set a grid of discrete hyperparameter values.
- Set a metric for scoring model performance.
- Search exhaustively through the grid.
- For each set of hyperparameters, evaluate each model's CV score.
- The optimal hyperparameters are those of the model achieving the best CV score.

Grid search cross validation: example

- Hyperparameters grids:
 - `max_depth = {2,3,4}`,
 - `min_samples_leaf = {0.05, 0.1}`
- hyperparameter space = $\{ (2,0.05) , (2,0.1) , (3,0.05), \dots \}$
- CV scores = $\{ \textit{score}_{(2,0.05)} , \dots \}$
- optimal hyperparameters = set of hyperparameters corresponding to the best CV score.

Inspecting the hyperparameters of a CART in sklearn

```
# Import DecisionTreeClassifier  
from sklearn.tree import DecisionTreeClassifier  
  
# Set seed to 1 for reproducibility SEED = 1  
  
# Instantiate a DecisionTreeClassifier 'dt'  
dt = DecisionTreeClassifier(random_state=SEED)
```


Inspecting the hyperparameters of a CART in sklearn

```
# Print out 'dt's hyperparameters
```

```
print(dt.get_params())
```

```
{'class_weight': None, 'criterion': 'gini',  
 'max_depth': None, 'max_features': None,  
 'max_leaf_nodes': None,  
 'min_impurity_decrease': 0.0,  
 'min_impurity_split': None,  
 'min_samples_leaf': 1,  
 'min_samples_split': 2,  
 'min_weight_fraction_leaf': 0.0, 'presort': False,  
 'random_state': 1,  
 'splitter': 'best'}
```

```

# Import GridSearchCV
from sklearn.model_selection import GridSearchCV # Define the grid
of hyperparameters 'params_dt' params_dt = {
    'max_depth': [3, 4, 5, 6],
    'min_samples_leaf': [0.04, 0.06, 0.08],
    'max_features': [0.2, 0.4, 0.6, 0.8]
}
# Instantiate a 10-fold CV grid search object 'grid_dt' grid_dt =
GridSearchCV(estimator=dt,
              param_grid=params_dt,
              scoring='accuracy', cv=10,
              n_jobs=-1)
# Fit 'grid_dt' to the training data grid_dt.fit(X_train,
y_train)

```

Extracting the best hyperparameters

```
# Extract best hyperparameters from 'grid_dt' best_hyperparams =  
grid_dt.best_params_  
print('Best hyperparameters:\n',  
best_hyperparams)
```

Best hyperparameters:

```
{'max_depth': 3, 'max_features': 0.4, 'min_samples_leaf': 0.06}
```

```
# Extract best CV score from 'grid_dt' best_CV_score =  
grid_dt.best_score_  
print('Best CV accuracy'.format(best_CV_score))
```

Extracting the best estimator

```
# Extract best model from 'grid_dt' best_model =  
grid_dt.best_estimator_  
  
# Evaluate test set accuracy  
test_acc = best_model.score(X_test,y_test)  
  
# Print test set accuracy  
print("Test set accuracy of best model: {:.3f}".format(test_acc))
```

Test set accuracy of best model: 0.947

Tuning an RF's Hyperparameters

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Random Forest (RF) is an ensemble learning algorithm based on **bagging (Bootstrap Aggregating)** that combines multiple decision trees to improve prediction accuracy and robustness. It can be used for both **classification** and **regression** tasks and is one of the most widely used machine learning algorithms due to its simplicity, effectiveness, and ability to handle various types of data.

Random Forests Hyperparameters

- CART hyperparameters
- number of estimators
- bootstrap
-

Tuning is expensive

Hyperparameter tuning:

- computationally expensive,
- sometimes leads to very slight improvement,

Weight the impact of tuning on the whole project.

Inspecting RF Hyperparameters in sklearn

```
# Import RandomForestRegressor
from sklearn.ensemble import RandomForestRegressor

# Set seed for reproducibility SEED = 1

# Instantiate a random forests regressor 'rf' rf =
RandomForestRegressor(random_state= SEED)
```



```
# Inspect rf's hyperparameters
```

```
rf.get_params()
```

```
{'bootstrap': True, 'criterion': 'mse',  
 'max_depth': None, 'max_features': 'auto',  
 'max_leaf_nodes': None,  
 'min_impurity_decrease': 0.0,  
 'min_impurity_split': None,  
 'min_samples_leaf': 1,  
 'min_samples_split': 2,  
 'min_weight_fraction_leaf': 0.0,  
 'n_estimators': 10,  
 'n_jobs': -1, 'oob_score':  
 False, 'random_state': 1,  
 'verbose': 0, 'warm_start':  
 False}
```

Basic imports

```
from sklearn.metrics import mean_squared_error as MSE
```

```
from sklearn.model_selection import GridSearchCV #
```

Define a grid of hyperparameter 'params_rf' params_rf = {

```
    'n_estimators': [300, 400, 500],
```

```
    'max_depth': [4, 6, 8],
```

```
    'min_samples_leaf': [0.1, 0.2],
```

```
    'max_features': ['log2', 'sqrt']
```

```
}
```

Instantiate 'grid_rf'

```
grid_rf = GridSearchCV(estimator=rf,
```

```
                        param_grid=params_rf, cv=3,
```

```
                        scoring='neg_mean_squared_error', verbose=1,
```

```
                        n_jobs=-1)
```

Searching for the best hyperparameters

```
# Fit 'grid_rf' to the training set grid_rf.fit(X_train,  
y_train)
```

Fitting 3 folds for each of 36 candidates, totalling 108 fits

[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 10.0s

[Parallel(n_jobs=-1)]: Done 108 out of 108 | elapsed: 24.3s finished

RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=4, max_features='log2',
max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=0.1, min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=400, n_jobs=1, oob_score=False, random_state=1, verbose=0,
warm_start=False)

Extracting the best hyperparameters

```
# Extract the best hyperparameters from 'grid_rf' best_hyperparams =  
grid_rf.best_params_
```

```
print('Best hyperparameters:\n', best_hyperparams) Best
```

hyperparameters:

```
{'max_depth': 4,  
 'max_features': 'log2',  
 'min_samples_leaf': 0.1,  
 'n_estimators': 400}
```

Evaluating the best model performance

```
# Extract the best model from 'grid_rf' best_model =  
grid_rf.best_estimator_  
  
# Predict the test set labels y_pred =  
best_model.predict(X_test) # Evaluate the test set  
RMSE  
  
rmse_test = MSE(y_test, y_pred)**(1/2) # Print the test  
set RMSE  
  
print('Test set RMSE of rf: {:.2f}'.format(rmse_test))
```

Test set RMSE of rf: 3.89